

Predicting FOMC Actions Using NLP; Assessing the Impact of Matrix Sparsity and Regularization

Jeremy Lao - jjl359
NYU Computer Science

John Reynolds - jr4716
NYU Computer Science

May 14, 2019

Abstract

The Federal Open Market Committee (FOMC) meets throughout the year to set the target Federal Funds rate. The Federal Funds rate is one of the primary monetary policy tools, and it impacts the cost of borrowing money globally. This paper explores ML and NLP methods, examines the impact of matrix sparsity from NLP methods and regularization from ML models, to predict the potential outcome of an upcoming FOMC using past FOMC statements and meeting minutes to train ML and NLP models.

1 Introduction

In this paper we use Natural Language Processing (NLP) and Machine Learning (ML) methods to predict Federal Open Market Committee (FOMC) rate actions (hold or change) using text from FOMC meeting minutes, Board of Governors speeches, and FOMC post-meeting statements.

For this analysis we used 713 separate documents (roughly 1.6 Million Words) and 148 Fed Decisions for our data set. We employ the NLP bag-of-words model for our analysis. Bag-of-words models have high dimensionality, meaning the space is extremely sparse with a large amount of features leading to a sparse feature matrix. We found in our analysis that some of the feature matrices contained over two billion elements with as little as 0.25% of the elements containing a non-zero value.

Due to these modeling challenges, the analysis and techniques we used address the issue of features being magnitudes larger than the label size ($p \gg n$), and the extreme sparsity in the feature matrix. In order to determine the appropriate model, n-gram size, and model parameters, we created a simulator to detect positive or negative sentiment from randomly generated text.

Our model simulator found that as the number of sentiment words increased, the prediction accuracy approached close to 1.0 with high F1 scores. For the study the analysis hones in on the mix of distributions containing hawkish or dovish (positive-action) words.

1.1 git repository for project

<https://www.github.com/jrrpanix/ML9>

2 Objectives

1. Given the proliferation of robust open-source ML and NLP packages, we were motivated to understand the characteristics of the feature matrix output (i.e., document-term matrix)

from the text processing functions (i.e., CountVectorizer) analyzing the number of features, matrix sparsity, and matrix norm generated from utilizing package functions.

2. Assess the effectiveness of Multinomial Naive Bayes and Logistic Lasso applied to problems where the feature size is orders of magnitudes larger than the label size ($p \gg n$).
3. Determine whether ML and NLP methods can be used to predict FOMC action (no change or a change in the base rate) using FOMC statements, minutes, and Board of Governors speeches.
4. Measure model performance on F1 scores because 70% of the actions over the nineteen year period we are measuring are “no action”. F1 scores ensure that the accuracy of the model is not based purely on a naive “no action” guess, which in the best case would give a 70% model accuracy.
5. Create scraping and cleaning tools to assemble a corpus of documents containing Speeches, Statements, Minutes and Actions from the Federal Reserve.

3 Methodology

3.1 Data

The time frame of our study covered 2000 to 2018. We primarily utilized three sources for our data, FOMC statements, FOMC meeting minutes, and Board of Governors speeches. Details on data-sources, extraction and scrubbing are left to the appendix to focus this paper on analyzing our methods and results. Preparing the data was a large amount of work and we are not aware of any data sets that combine the three types of documents and mapped the documents to an FOMC Action.

3.2 Feature Modeling and Data Label Mapping

Documents are mapped to the first FOMC rate decision on or before the document date at time t_i , where j = Statement or j = Speech or j = Minutes

$$\text{Document}_{j,t_i} \rightarrow \text{Decision}_{t_i+k} = \begin{cases} 1, & \text{if Decision} = \text{change}, \\ 0, & \text{if Decision} = \text{no change} \end{cases}$$

Model 1: Unstacked

The **Unstacked** model maps each document to a feature vector generated by CountVector given an n-gram (n_1, n_2) where $n_1 \leq n_2$ and n_i are integers > 0 . Given that our corpus consisted of approximately 713 documents, this will generate a feature matrix with 713 rows and K columns where each column will be a unique instance of the word sequence generated by the n-gram.

$$E[D_{t_i}] = f(\text{ngram}(n_1, n_2)[\text{Document}_{j,t_i-k}])$$

Model 2: Stacked

The **Stacked** model maps the set of documents which occur after the previous decision date but on or before the next decision date to one rate decision. We effectively concatenated the documents between decisions into one large document. The “stacked” model reduces the number of documents from 713 to 148, and we expect this method to reduce the dimensionality of the space. This model maintains the same n-gram based feature vector where $n_1 \leq n_2$ and n_i are integers > 0 .

$$E[D_{t_i}] = f(\text{ngram}(n_1, n_2)[\text{Document}_{j_1,t_i-k_1}, \text{Document}_{j_2,t_i-k_2}, \dots]), t_i - 1 < t_i - k_i < t_i + 1$$

3.3 Matrix Sizes and Sparsity from CountVectorizer

For this study we used term frequencies generated from **CountVectorizer** (**Tfidf** can also be used, but it did not alter the conclusions of this study). When an n-gram is applied to a corpus of documents D , one can deduce from basic combinatorics that a large number of features will be generated from various mixtures of case studies.

As we mentioned in our introduction, the classic problem with NLP bag-of-words models is that most of the column entries (n-grams) are zero, resulting in extremely sparse matrices. Sparsity is defined as:

$$\text{sparsity} = \frac{\text{non-zero elements}}{\text{total elements}}$$

The table below gives a measure of sparsity and the resulting number of elements in each n-gram generated matrix (units are in billions). We empirically observed that the “stacked” model reduces the sparsity by a factor of four and the matrix size by a factor of five. The L1 matrix norm for all of the matrices ranged between 1200 to 1600.

CountVector Matrix Size				
	Unstacked		Stacked	
N-gram	sparsity	size(billions)	sparsity	size(billions)
1:1	0.03628801	0.008561	0.10635940	0.001757
1:2	0.00544312	0.170058	0.02081416	0.036107
1:3	0.00368457	0.439898	0.01508941	0.092869
3:5	0.00238128	0.900721	0.01059293	0.173040
4:7	0.00224699	1.332560	0.00996789	0.271117
5:10	0.00218455	1.995502	0.00972276	0.402041
8:10	0.00213633	1.019811	0.00955894	0.198627
10:15	0.00210643	2.097187	0.00945586	0.398312
20:20	0.00205055	0.347043	0.00931157	0.074799

3.4 Apply ML to Sparse Matrices generated from CountVectorizer

Based upon research into classification models for sparse data, we chose to test Naive Bayes and Logistic Regression with L1 regularization.

- Naive Bayes

$$P(A|x_1, x_2, \dots, x_n) = p(A) \prod_{i=1}^n p(x_i|A)$$

- Logistic Lasso

$$A = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i} + L1$$

For the “Unstacked” document case, our research into classification model performance on sparse data found that Naive Bayes classification models performed particularly well when encountering cases of missing data. Naive Bayes takes advantage of the assumption that the attributes are independent and ignores missing attributes when calculating the probability. Therefore, Naive Bayes is generally considered a fairly decent performer with greater sparsity.

Based on the table above called “Count Vector Size”, we decided to test Logistic Lasso as we found in our research that Logistic Lasso can handle sparsity through identifying a set of relevant variables out of the large collection of candidates making it suitable for discovering interactions in regression problems (HW1 problem 2).

3.4.1 Training, Testing, and Determining the Efficacy of the Models

For this study we randomly sample 75% of the corpus for training and then predict Fed action using the remaining 25% of the data and calculate

- Accuracy
- Precision
- Recall
- F1 Score

Approximately 70% of the time there is no Action so a model that does nothing will have an accuracy of 70%, but a recall of 0. The results are summarized using F1 score.

3.4.2 The impact of regularization coefficient α on logistic regression

For **Logistic Lasso** we are interested in examining the impact of the regularization parameter α on the fraction of non-zero coefficients remaining after regularization is applied and on the impact of the predictive score on test data. The results show the regularization parameter at the extremes produces the worst results and the stacked model has a higher fraction of coefficients eliminated from logistic lasso. The expected result is the larger regularization parameter reduces the fraction of non-zero coefficients.

Comparison of Logistic Lasso varying regularization coefficient α					
		Unstacked		Stacked	
N-gram	α	F1	frac remain	F1	frac remain
10:15	3.333	0.454	0.000040	0.720	0.000016
10:15	2.000	0.575	0.000066	0.897	0.000029
10:15	1.000	0.647	0.000148	0.893	0.000055
10:15	0.500	0.679	0.000291	0.913	0.000094
10:15	0.200	0.708	0.000339	0.913	0.000134
10:15	0.100	0.710	0.000463	0.923	0.000161
10:15	0.050	0.711	0.000619	0.925	0.000206
10:15	0.020	0.715	0.000764	0.923	0.000243
10:15	0.010	0.713	0.001045	0.913	0.000315
10:15	0.002	0.730	0.001659	0.906	0.000517
10:15	0.000	0.669	0.978049	0.669	1.000000

3.4.3 A comparison of Multinomial Naive Bayes to Logistic Regression with Extremely Sparse Feature Matrices

Under extreme levels of sparseness (sparseness measure of less than 0.0021) from feature matrices generated by CountVectorizer, Logistic Regression and Logistic Lasso break down with greater sparseness while Multinomial Naive Bayes can still produce usable results.

However, we found with less matrix sparsity and a lower Logistic Lasso model penalty (i.e., $\alpha = 0.002$), the Logistic Lasso model performance improves. Based on F1 scores we observed that Logistic Lasso with low penalties performs as well or better than Naive Bayes when sparseness exceeded 0.0025.

Logistic Lasso with a larger penalty (i.e., > 2.0) does not perform well with extremely sparse matrices.

Deterioration of Logistic Lasso with sparse matrices relative to Naive Bayes				
N-gram	sparsity	F1 Score		
		Naive Bayes	Logistic Lasso	
		smooth= 0.0	$\alpha = 0.002$	$\alpha = 1.000$
60	0.001979	0.647	0.463	0.256
50	0.002018	0.619	0.565	0.342
30	0.002038	0.724	0.620	0.460
25	0.002063	0.720	0.664	0.490
20	0.002067	0.725	0.678	0.520
12	0.002069	0.682	0.743	0.656

Below is a plot showing Naive Bayes giving consistent results relative to Logistic Lasso .




Figure 1: F1 performance of LL and NB - Unstacked Model

3.4.4 Impact of N-gram sequences on results

We varied the N-gram sequences to observe its impact on our model predictions. Based on our empirical observations there is a definite range of N-gram sequences where all of the models greatly improve (based on F1 scores). All of the models showed great improvement when the N-gram sequences ranged between 4:6 and 10:15. As the sequences (N-grams) approached 20 words and larger, the results (model performance) started to deteriorate as the feature matrices became more sparse. Below is a graph of Logistic Lasso and Naive Bayes plotted with different regularization coefficients and varying N-gram sequence.

The interesting part about our findings is that Federal Reserve commentary tends to have sequences of hawkish or dovish sentiment (FOMC action) which can't be expressed in a few words. This expression of hawkishness or dovishness in larger N-grams may be the reason why smaller B-gram sequences did not perform as well.

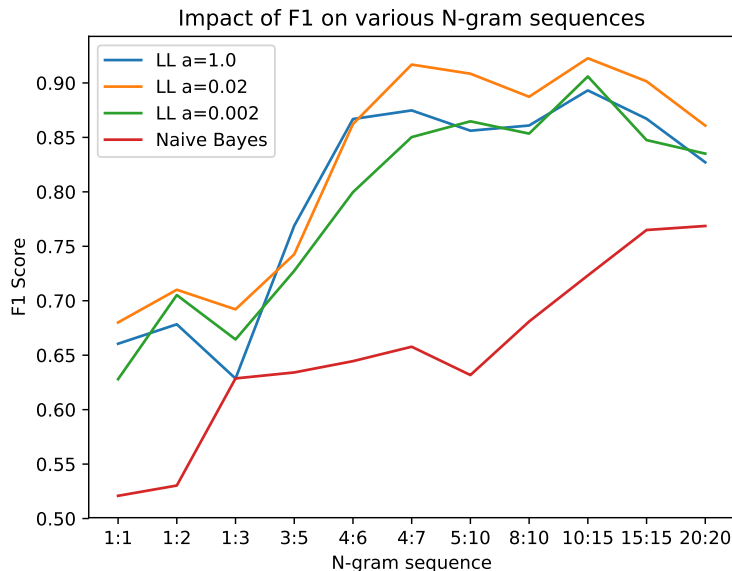


Figure 2: N-gram sequence LL and NB - Stacked Model

4 Conclusion

Reducing the sparsity of the matrices generated from CountVectorizer by stacking the documents greatly improved the results.

We were able to attain maximum test accuracy of 95.9% using a 10 : 15 n-gram, producing a recall of 93.4% and an F1 score of 92.5%. Because of the large number (i.e., millions) of features in the feature matrix, Logistic Lasso with small penalty proved to be the most effective method for reducing the number of coefficients, leading to improved F1 scores.

The fraction of non-zero coefficients remaining after applying Logistic Lasso ranged between $3e - 05$ to $1.6e - 03$. Logistic Lasso was very effective in dealing with feature matrix sparsity, and based on what we learned from HW1 question two is the appropriate model choice for our classification problem. In order to illustrate the effectiveness, we measured the fraction of non-zero coefficients using a Logistic Regression model and the feature matrix had $\approx 97\%$ non-zero coefficients.

Logistic Lasso (and Logistic Regression generally) broke down at the most extreme levels of sparsity < 0.0020 . In general, we found that setting the regularization parameters at the extremes ($\alpha = 0.0$ or $\alpha > 2.0$) produced inferior results.

We found that n-gram sequences of varying between 4:7 to 10:15 provided the best results. Smaller n-gram sequences (< 4 words) and larger sequences (> 20 do to extreme sparsity) did not generate good results.

Whereas Multinomial Naive Bayes with Laplace smoothing = 0 was the most stable method for producing useful results at the most extreme measures of matrix sparsity. In general we found that when using Multinomial Naive Bayes setting the parameters at no smoothing or large smoothing proved to give the best results.

Overall our research found that Logistic Lasso with smaller penalty $\alpha < 0.02$ produced the most stable results for predicting FOMC rate decisions from FOMC statements, minutes, and Board of Governors speeches.

The limitation in this study is the size of the data sets. While we were able to identify characteristics that would help predictive power the strength of the conclusions are tempered by the amount of data on hand.

A Appendix

A.1 What is the FOMC

The FOMC sets the Federal Funds target range (target rate prior to 2009). The permanent members are the Board of Governors of the Federal Reserve System, President of the Federal Reserve Bank of New York, and the rest of the seats are rotated through the Presidents of the other Federal Reserve Banks. There are twelve Reserve Banks and the Board of Governors oversees the activity, operations, and policies of the Federal Reserve System.

A.2 Why Does the Federal Funds Target Range Matter?

The Federal Funds Target Range is set by the FOMC. The Federal Funds Target Range is the range where the Federal Funds Effective rate (calculated as the volume weighted median of eligible transactions reported on FR 2420) is expected to fix on a daily basis. The Federal Funds rate is the primary monetary policy instrument of the Federal Reserve System, and it influences the level of interest rates domestically and globally. For example, if the Federal Funds effective rate were 5 percent, then the interest rate on a 30 year mortgage would be 5 percent or greater.

A.3 Data

A.3.1 Data Sources

The time frame of our study covered 2000 to 2018. The primary source of our textual data were FOMC statements, FOMC meeting minutes, and Board of Governors speeches from <https://www.federalreserve.gov>. The data for the Federal Funds target rate and target range (post 2009) is from FRED St. Louis. FRED offers a wealth of economic data and information to promote economic education and enhance economic research. FRED is updated regularly and allows 24/7 access to regional and national financial and economic data.

A.3.2 Scraping and Text Pre-Processing

We collected data from the Federal Reserve website and pre-processed the data for the document-term matrix.

We utilized Python packages *beautifulsoup4*, *re*, and *urllib* to scrape contextual data from the Federal Reserve website. Our scraping algorithm used regular expressions to handle and remove extraneous *html* and *javascript* text collected by the page scraper. We had to handle non UTF-8 characters upfront and remove them at this stage altogether. We collected over six hundred documents from the Federal Reserve website.

After scraping the data, we pre-processed the text by removing all punctuation, ensuring proper spacing between words, setting all words to lower case, and making all numbers *d* to reduce the dimensionality document-term matrix. We also used Regex to detect direct references to the Federal Funds target rate/range and transformed those references (a mixture of numbers, punctuation, and the word 'percent') to *ddpercentrate*. Since we were modeling the sentiment of FOMC rate action, this was one of our strategies to directly capture the target rate in the document-term matrix.

This formed the bulk of our text pre-processing and it allowed us to essentially utilize the stored text in the document-term matrix for NLP and ML models.

A.3.3 Meta-Data

Below is a table of the FOMC document metadata:

FOMC Documents					
Document Type	Years	Number Documents	Total Words	Avg Words Per Document	StDev
Minutes	2000-2018	149	738,656	4,957	1,284
Statements	2000-2018	160	62,350	389	185
Speeches	2011-2018	403	915,359	2,271	1,539

Statements are post-meeting communiqué and are generally short, but contains a summary of the committee’s reasoning for the decision. Minutes are longer and contain the Staff and Committee member current economic assessments and outlooks. The speeches in the corpus are from the Board of Governors and may contain clues of their thoughts on their current economic outlook.

A.4 Latent Dirichlet Allocation

We explored latent topic modeling in our research to determine how a machine learning algorithm would classify $k = 3$ topics, a proxy for the Federal Reserve’s three objectives.

Given the Federal Reserve’s mandates of stable prices, maximum employment, and financial stability, we explored Latent Dirichlet Distribution of the speeches and minutes’ word distribution over three latent topics. We were interested in exploring how well a topic model would distribute words over the three objectives and if the meaning of the distributions were relatively clear.

```
Topic #0: inflation|economic|dd|prices|growth|quarter|participants|remained|market|consumer|continued|year|rate|spending|recent
Topic #1: financial|dd|banks|capital|federal|reserve|bank|crisis|firms|federal reserve|risk|important|banking|community|market
Topic #2: rate|inflation|dd|policy|market|economic|federal|percenttarget|labor|monetary|growth|participants|term|monetary policy|longer
```

Figure 3: Top 15 Words Distributed Over Three Latent Topics

We found that the topic model generally distributed the speeches and minutes into the three categories fairly well. It was interesting that "Topic 1" referred solely to financial stability, "Topic 0" primarily referred to inflation and spending, and "Topic 2" returned a word distribution that included labor, growth, and lower for longer monetary policy.

References

- [1] Hansen,Stephen, McMachon, Michael *Transparency and Deliberation with the FOMC:a Computational Linguistics Approach*, CEPR Discussion Paper, (2014)
- [2] Puri, Indira *Using Machine Learning to Predict Interest Rate Changes from Federal Reserve Proceedings*
- [3] Ranganath, Rajesh *NYU - CSCI-GA.2565 ML Lecture Slides Spring 2019*
- [4] Li, Xiang et al. *The Convergence Behavior of Naive Bayes on Large Sparse Datasets*, ACM Vol 11 Issue 1, August 2016
- [5] Zou, Yuan and Roos, Teemu *Sparse Logistic Regression with Logical Features*, Helsinki Institute for Information Technology (HIIT), Springer International, 2016